

Agree on the Evidence: Grounded Cross-Model Agreement for Trustworthy, Explainable LLM Phishing Detection

Abstract

Large language models (LLMs) now flag phishing websites with high accuracy, but a deployed detector consumes a *label* and discards the *reason* – and an action such as a block or a take-down needs a reason an operator can check. Worse, model rationales are often produced after the fact: a panel of detectors can agree on the right verdict for the wrong, or even fabricated, reason, a failure that label confidence and majority voting cannot see. We propose PHISHPROOF, a multi-model detector that decides how much to trust a verdict from how much independent models agree on the underlying evidence. Each model in a small, diverse panel returns its verdict together with a short list of typed, checkable evidence units; we score agreement at the level of that evidence, weight it by whether each cited unit actually verifies against the page, and act only when this grounded agreement is high – returning the agreed evidence as the explanation and abstaining otherwise. Phishing is the natural home for the idea: its evidence is objectively checkable, so shared hallucinations are caught, and a delayed but high-precision oracle (later blocklisting) calibrates the trust score. Framed as selective prediction, PHISHPROOF improves the risk-coverage trade-off over single-model confidence, self-consistency, and label voting across a six-month benchmark, and flags the right-label, wrong-reason pages the baselines act on over-confidently.

CCS Concepts

• Security and privacy → Phishing; • Computing methodologies → Ensemble methods; Uncertainty quantification.

Keywords

phishing detection, large language models, selective prediction, multi-model agreement, explanation faithfulness, abstention, trustworthy ML

1 Introduction

A phishing website is a consequential object: when a detector marks a page as phishing, the verdict drives a blacklist entry, a browser warning, or a take-down request that removes a live service from the web. Large language models have raised the accuracy of such detectors sharply, reading a page’s URL, DOM, rendered screenshot, and certificate to decide whether it impersonates a brand it does not own [13, 15]. Yet accuracy is only half of what a deployment needs. The analyst who acts on a verdict, the operator who must justify a take-down, and the user who sees a warning all need a *reason* they can check – and they need to know *when* the detector is trustworthy enough to act on automatically. For a decision that can be contested, the bottleneck is no longer whether the model is usually right, but whether a particular verdict can be *trusted and explained*.

Existing detectors approach this only indirectly. One line runs a single LLM over the page and reports a class probability or a verbalized confidence [15]; reliability is then a property of that one




Figure 1: PHISHPROOF in one picture. A diverse panel labels a page and *cites checkable evidence*. We trust the verdict only when the panel agrees on evidence that *verifies against the page*, return that evidence as the explanation, and *abstain on right-label, wrong-reason* pages where the models agree on the label but not on why.

model’s calibration [7]. A second line aggregates several models or several samples – self-consistency over one model [24], majority voting over an ensemble [22], or weighted label agreement in the style of crowd-sourcing and label-free evaluation [4, 5] – and reads reliability off how often the votes concur. A third, recent line builds multi-agent detectors whose agents specialize on different signals and whose rationales are merged into a single explanation [12, 25]. Together these families cover much of the space of accuracy and aggregation.

They share a blind spot, visible in a routine case (Fig. 1). A page impersonating a bank is shown to a panel of three diverse models; all three return “phishing,” so every aggregation method above reports high confidence. But the models reached that verdict for different reasons: one cited the morphed logo, one cited a credential form posting off-domain, and one cited a certificate mismatch that does not actually hold on the page. First, the detectors emit a label and discard these reasons, so nothing auditable reaches the operator (L1). Second, reliability is read at the *label*, where the three agree, and is therefore blind to the fact that they *disagree on the evidence* – the verdict is right for the wrong reason (L2). Third, where the models do agree on a reason, that agreement can be a shared hallucination of evidence that is not on the page (L3). Fourth, even a detector that wanted to gate on agreement has no ground-truth at inference time to calibrate how much agreement is enough, and the threshold it picks drifts as campaigns evolve (L4). A page that is right for the wrong reason is not a curiosity: it is exactly the page a small change to the cited cue would flip, and exactly the page an operator should not act on blindly.

Two properties of phishing, not shared by generic question answering, make a direct remedy possible now. The evidence a detector relies on is *objectively checkable*: a form-action domain, a certificate organization, a redirect target, and a rendered logo can each be re-derived from the page itself. And phishing carries a delayed but high-precision *oracle* – a malicious URL is later blocklisted or taken down [18, 20] – that supplies real labels after the fact. Modern LLMs also emit typed, structured outputs cheaply

enough to ask every model for its evidence, not just its label. The opportunity is to stop treating reliability as a property of labels and start treating it as a property of *evidence that several independent models share and that the page confirms*.

To this end, we present PHISHPROOF, a multi-model phishing detector that trusts a verdict in proportion to grounded cross-model agreement on the evidence. A small, deliberately diverse panel labels the page and each model cites a short list of typed evidence units from a fixed schema. PHISHPROOF measures agreement at the level of those units rather than the label, weights it by a verifier that checks whether each cited unit holds on the page, and combines the two into a single Grounded Evidence-Agreement score. It acts only when that score is high – returning the agreed, verified units as a human-readable explanation – and abstains otherwise, escalating the page to deeper analysis or a human. A delayed-oracle calibrator maps the score to a probability of correctness so the operating point is set against real outcomes rather than guessed. The novelty over prior multi-model detectors is the unit of agreement: not the label, and not one model’s introspective faithfulness, but evidence that is simultaneously cross-model-agreed and page-verified.

Realizing this raises four difficulties that the design must rule out. Agreement over structured evidence must be defined so that it is not label agreement in disguise – two models can name the same brand while resting on different cues. Shared hallucination must be caught cheaply, without trusting the models to police themselves. The trust score must be calibrated without any inference-time label and must degrade gracefully as campaigns drift, using only the oracle’s late arrivals. And the panel must stay small, since every added model is a forward pass, so the gain has to come from *diversity* rather than count.

We make the following contributions.

- **Problem framing.** We identify the *right-label, wrong-reason* failure of multi-model phishing detectors and cast trustworthy detection as selective prediction over an evidence-agreement signal (§3).
- **Method.** We introduce PHISHPROOF and its Grounded Evidence-Agreement score, which measures cross-model agreement at the level of typed, page-verified evidence and turns it, via delayed-oracle calibration, into a selective act/abstain decision with the agreed evidence as the explanation (§4).
- **Evidence.** On a six-month phishing benchmark, PHISHPROOF improves the risk-coverage trade-off over single-model confidence, self-consistency, and label voting, and abstains on the fragile right-label, wrong-reason pages that the baselines act on (§6).
- **Analysis.** We isolate the contribution of evidence-level agreement, groundedness, panel diversity, and calibration, and report behavior under temporal drift (§6).

The rest of the paper is organized as follows. §2 positions PHISHPROOF against label-level reliability and multi-agent detection. §3 formalizes grounded evidence-agreement and the selective objective. §4 details the panel, the score, and the calibrated decision rule. §5 discusses deployment considerations, and §6 reports the evaluation.

2 Related Work

LLM and reference-based phishing detection. The dominant modern detectors decide whether a page impersonates a brand it does not own. Reference-based systems match a rendered page against a brand library of logos and layouts [1, 14, 16], and recent work refreshes that reference set online [17] or replaces it with an LLM that recognizes the brand and checks the domain directly [13, 15]. A parallel line treats the page as URL or image features for a learned classifier [26], surveyed broadly [21]. These systems optimize and report *aggregate* accuracy. Departing from this line, PHISHPROOF adds no new base detector and makes no accuracy claim; it is a layer that decides, *per verdict*, whether the panel’s decision can be trusted and acted on (§4.4).

Multi-agent and ensemble phishing detection. Closest to us are detectors that run several models. Majority voting over LLMs has been studied directly for phishing [22]; multi-agent systems assign agents to different signals and either debate to a verdict or merge their rationales into one explanation [12, 25], sometimes backed by an episodic memory of past cases [3]. Unlike these systems, which aggregate or consolidate at the level of the *label* or a fused rationale, PHISHPROOF measures agreement at the level of individually cited evidence units ((2)) and treats *disagreement on evidence* as the signal to abstain rather than a conflict to resolve.

Label-level reliability and label-free evaluation. A verdict’s reliability is usually read from the label distribution: post-hoc confidence calibration [7], agreement across samples of one model [24], or agreement across many models, as in crowd-sourcing aggregation [4], agreement-on-the-line under shift [2], and trust calibration for phishing detectors [5]. Where these estimate reliability from *how often labels concur*, PHISHPROOF estimates it from *whether the reasons concur and hold*, which separates pages on which the panel agrees by coincidence from those on which it agrees because the page truly carries the evidence ((4)).

Explanation faithfulness. Causal audits show that an LLM’s stated rationale is often not the driver of its decision [9], and multi-agent debate has been used to *improve* the faithfulness of generated explanations [10]. Such audits operate on a single model and require counterfactual re-execution. In contrast, PHISHPROOF uses cross-model agreement on grounded evidence as an inexpensive faithfulness *proxy*: a reason that several independent models cite and that the page confirms is, by construction, less likely to be a post-hoc rationalization, at no extra re-execution cost.

Selective prediction. We build on selective classification, which trades coverage for accuracy via an abstention rule [6], and on distribution-free uncertainty quantification [23]. Building on this foundation, PHISHPROOF supplies the missing ingredient for the phishing setting: an abstention score derived from grounded cross-model agreement, and a calibration target – the delayed blocklisting oracle – that lets the risk-coverage operating point be set against real outcomes rather than a held-out guess.

Position and scope. PHISHPROOF contributes a single new object, the Grounded Evidence-Agreement score, and the selective detector built on it. ?? positions it against one representative per family along five capabilities – evidence-level agreement, page-grounding, oracle calibration, selective abstention, and a verifiable explanation

as output – and PHISHPROOF is the only entry that supports all five. We focus on per-verdict trust and explanation for a multi-model phishing detector; the framework is agnostic to the specific base models and evidence schema. We do *not* claim (i) a new base detector or higher full-coverage accuracy – PHISHPROOF layers over any panel; (ii) robustness to an adversary who controls the panel models or corrupts the verifier’s inputs, a harness-level threat we treat as out of scope; or (iii) coverage guarantees on fully cloaked, novel pages for which the oracle has not yet arrived, where the detector necessarily abstains rather than errs.

3 Problem Formulation

Setting. A detector receives a webpage x , represented by its URL, DOM, rendered screenshot, TLS certificate, and redirect chain, and must decide whether x is phishing. We assume access to a panel $\mathcal{M} = \{m_1, \dots, m_M\}$ of M diverse LLM detectors (different model families, including at least one vision-language model that reads the screenshot), and to a fixed, typed *evidence schema* $\mathcal{U}$. A unit $u = (t, v) \in \mathcal{U}$ pairs a type t (e.g. form-action-domain, cert-org, redirect-target, logo-brand, credential-intent) with a value v . Notation is summarized in Tab. 1.

Model outputs. Run on x , each model returns a verdict and a short list of cited units:

$$m_i(x) = (y_i(x), E_i(x)), \quad y_i \in \{\text{phish}, \text{benign}\}, \quad E_i \subseteq \mathcal{U}. \quad (1)$$

The panel label is the majority vote $\hat{y}(x) = \text{maj}_i y_i(x)$. A *consensus* unit is one a strict majority of the panel cites with a consistent value, $C(x) = \{u : |\{i : u \in E_i\}| > M/2\}$.

Prior formulation and its limits. Existing detectors estimate a per-verdict reliability $r(x)$ from the *label* channel alone – a single model’s confidence, or the concordance of $\{y_i\}$ [4, 7, 22]. This label-centric formulation has four limitations, which we restate as the requirements PHISHPROOF must meet (Tab. 2). It returns no checkable artifact, only $\hat{y}$, so an action cannot be justified (**L1**; **R1**: emit structured evidence). It is a function of $\{y_i\}$ only, so it cannot distinguish a panel that agrees *because* the page carries the evidence from one that agrees on the label while resting on different cues (**L2**; **R2**: score agreement on evidence). It treats a cited reason as true, so a reason all models hallucinate is rewarded (**L3**; **R3**: ground evidence to the page). And it offers no target against which to set its operating point as campaigns drift (**L4**; **R4**: calibrate against the delayed oracle).

Grounded evidence-agreement. We replace label reliability with a score on the *evidence* channel. Define the evidence-level agreement and the groundedness of the consensus,

$$A(x) = \frac{|C(x)|}{|\bigcup_i E_i(x)|}, \quad (2)$$

$$G(x) = \frac{1}{|C(x)|} \sum_{u \in C(x)} g(u, x), \quad (3)$$

where $g(u, x) \in [0, 1]$ is a verifier that re-derives unit u from x (deterministic for form-action, cert-org, redirect; perceptual for logo-brand); $G(x)=0$ when $C(x)$ is empty. The *Grounded Evidence-Agreement* score is their product,

$$\text{GEA}(x) = A(x) \cdot G(x) \in [0, 1]. \quad (4)$$




Figure 2: PHISHPROOF pipeline. A diverse panel cites typed evidence units (C_1); we score cross-model agreement on those units (C_2 , **R2**) and their groundedness against the page (C_3 , **R3**); their product is the GEA score, which a delayed-oracle calibrator turns into a selective act/abstain decision with the agreed evidence as the explanation (C_4 , **R4**).

Table 1: Notation.

Symbol	Meaning
x	a webpage under test (URL, DOM, screenshot, cert, redirects)
$\mathcal{M} = \{m_1, \dots, m_M\}$	panel of M diverse detectors ($M=3$ default)
$\mathcal{U}$	fixed typed evidence-unit schema (§4.1)
$E_m \subseteq \mathcal{U}$	evidence units cited by model m on x
C	cross-model <i>consensus</i> units (cited by a majority)
$g(u, x) \in [0, 1]$	verifier: does unit u hold on x ?
$A(x) \in [0, 1]$	evidence-level agreement over $\{E_m\}$
$G(x) \in [0, 1]$	groundedness of the consensus units
$\text{GEA}(x) = A \cdot G$	Grounded Evidence-Agreement score
τ	abstention threshold (acts iff $\text{GEA} \geq \tau$)
$\hat{\rho} = \kappa(\text{GEA})$	calibrated map $\text{GEA} \rightarrow \mathbb{P}(\text{correct})$

Agreement on units the page does not confirm ($G \rightarrow 0$) and label agreement without evidence consensus ($A \rightarrow 0$) both drive GEA down, directly discharging **R2** and **R3**.

Objective. PHISHPROOF is a selective detector: it acts with $\hat{y}(x)$ when $\text{GEA}(x) \geq \tau$ and abstains ($\perp$) otherwise. For a threshold τ , *coverage* and *selective risk* are

$$\phi(\tau) = \mathbb{P}[\text{GEA} \geq \tau], \quad r(\tau) = \mathbb{P}[\hat{y} \neq y \mid \text{GEA} \geq \tau], \quad (5)$$

with y the true label. We seek a score whose risk-coverage curve $r(\phi)$ lies as low as possible (small area under it, AURC), and a calibrator $\kappa : \text{GEA} \mapsto \hat{\mathbb{P}}(\hat{y} = y)$ fit on the delayed oracle so that the operating point τ realizes a target risk (**R4**). *Outputs* per page: a verdict or abstention, the calibrated trust $\hat{\rho} = \kappa(\text{GEA})$, and, when acting, the consensus units $C(x)$ as the explanation.

Framework overview. Fig. 2 maps the requirements to four components evaluated in §6: evidence elicitation (C_1 , **R1**), evidence agreement (C_2 , **R2**), groundedness (C_3 , **R3**), and calibrated selection (C_4 , **R4**). §4 details each; Tab. 2 is the thread we follow through to the ablation.

Table 2: The traceable thread: each limitation (§1) becomes a requirement, a named component, and an ablation row (??).

Limitation	Requirement	Component	Ablation
L1 label-only, unverifiable output	R1 emit structured, checkable evidence	C_1 Evidence elicitation (§4.1)	output form
L2 reliability read at the label, blind to wrong reasons	R2 measure agreement at the evidence level	C_2 Evidence agreement A (§4.2)	–evidence
L3 shared hallucination of false evidence	R3 ground cited evidence to the page	C_3 Groundedness G (§4.3)	–grounded
L4 no ground-truth calibration; drift	R4 calibrate & select against the delayed oracle	C_4 Calibrated selection (§4.4)	–calib.; drift

4 The PHISHPROOF Detector

PHISHPROOF turns the score of (4) into a calibrated, selective verdict through the four components of Fig. 2. We describe each, then give the end-to-end procedure and its cost.

4.1 Evidence elicitation

Each model in the panel is prompted to return both a verdict and a short list of cited evidence units drawn from the fixed schema $\mathcal{U}$, using constrained decoding so that every cited unit is a well-typed (t, v) pair rather than free text (**R1**). The schema is small and auditable – six types covering the brand claim, the form-action domain, the certificate organization, the redirect target, credential intent, and the logo brand – and is the same for every model, so that their outputs are directly comparable. Panel *diversity* is a design choice, not an afterthought: the three models come from different families and at least one is a vision-language model that reads the screenshot, so that agreement reflects independent evidence rather than a shared prior (its value is isolated in §6).

4.2 Evidence agreement

To compare cited units across models we first *normalize values*: domains are reduced to their registrable eTLD+1, brands to a canonical lexicon entry, so that “paypal-secure.com” and “paypal-secure.com/login” count as the same form-action unit. The consensus $C(x)$ collects the value-normalized units a strict majority of the panel cites, and the evidence-level agreement $A(x)$ of (2) is the share of all cited evidence that falls in this consensus (**R2**). Agreement is thus high only when the panel converges on the *same* reasons, and is unaffected by the models concurring on the label alone.

4.3 Groundedness

A consensus can still be wrong if the models share a hallucination, so each consensus unit is checked against the page by a per-type verifier $g(u, x)$ (**R3**). The structural types are verified deterministically – the DOM is parsed for the form-action domain, the certificate for its organization, and the redirect chain for its final target – and return 0 or 1. The perceptual type (logo-brand) is verified by matching the rendered logo to the claimed brand and returns a similarity in $[0, 1]$. The groundedness $G(x)$ of (3) averages these over the consensus, and the score $\text{GEA} = A \cdot G$ of (4) collapses whenever the agreed evidence does not survive verification.

Algorithm 1 PHISHPROOF inference on a page x

```

1: for  $m_i \in \mathcal{M}$ :  $(y_i, E_i) \leftarrow m_i(x)$  ▷  $C_1$  elicit, (1)
2:  $\hat{y} \leftarrow \text{maj}_i y_i$ 
3:  $C \leftarrow \{u : u \text{ cited by } > M/2 \text{ models, value-normalized}\}$ 
4:  $A \leftarrow |C| / |\bigcup_i E_i|$  ▷  $C_2$ , (2)
5: for  $u \in C$ : compute  $g(u, x)$ ;  $G \leftarrow \text{mean}_{u \in C} g(u, x)$  ▷  $C_3$ , (3)
6:  $\text{GEA} \leftarrow A \cdot G$ ;  $\hat{\rho} \leftarrow \kappa(\text{GEA})$  ▷  $C_4$ , (4)
7: if  $\text{GEA} \geq \tau$  return  $(\hat{y}, \hat{\rho}, C)$ ; else return  $\perp$ 
```

4.4 Calibrated selection

Finally, PHISHPROOF turns the score into a decision (**R4**). Because phishing labels arrive late – a malicious URL is eventually block-listed or taken down [18, 20] – we fit an isotonic calibrator $\kappa : \text{GEA} \mapsto \hat{\mathbb{P}}(\hat{y} = y)$ on pages whose oracle label has already arrived, and refresh it on a sliding window so the mapping tracks drift. The detector acts with $\hat{y}(x)$ when $\text{GEA}(x) \geq \tau$, where τ is the smallest threshold meeting a target selective risk under κ , and abstains otherwise, escalating the page to deeper analysis or a human. When it acts, it returns the consensus units $C(x)$ as a verifiable explanation.

4.5 Procedure, guarantee, and cost

Walk-through. 1 runs the panel once (line 1), takes the majority label (line 2), and then does the work that distinguishes PHISHPROOF from label-level reliability: it forms the cross-model evidence consensus (line 3), scores agreement on that evidence rather than on the label (line 4), verifies each agreed unit against the page so a shared hallucination cannot inflate the score (line 5), and combines the two factors into GEA (line 6). The calibrated trust and the selective decision (lines 7) set the operating point against real outcomes. The novelty is concentrated in lines 3–6: the quantity that gates the verdict is grounded agreement on *evidence*, and the returned consensus is simultaneously the trust signal and the explanation.

Guarantee. Grounding is what rules out the shared-hallucination failure, and we can state this precisely.

ASSUMPTION 1 (VERIFIER SOUNDNESS). *For each unit type, $g(u, x) = 1$ if u holds on x and 0 otherwise (perceptual types: $g(u, x)$ is the true match score).*

PROPOSITION 1. *Under 1, any consensus unit that does not hold on x contributes 0 to $G(x)$; hence $\text{GEA}(x)$ is at most the grounded fraction of the consensus, and $\text{GEA}(x) = 0$ whenever the entire consensus is hallucinated.*

PROOF SKETCH. G is the mean of g over C ((3)); by 1 a non-holding unit adds 0 to that mean, so $G \leq |\{u \in C : u \text{ holds}\}|/|C|$, and $\text{GEA} = A \cdot G \leq G$ inherits the bound. If no consensus unit holds, $G = 0$ and thus $\text{GEA} = 0$. $\square$

The proposition is modest by design: it bounds, but does not eliminate, gauge error, since a verifier is sound for what it checks but a page can carry true-yet misleading evidence; we report verifier precision and recall against gold in §6.

Complexity. Per page, PHISHPROOF costs M panel forward passes plus $|C|$ verifier calls; the structural verifiers are $\mathcal{O}(1)$ DOM/certificate/redirect lookups and the single perceptual check is one embedding comparison, so the cost is dominated by the $M=3$ model calls and is

independent of the schema size beyond the units actually cited. Relative to a single-model detector this is a fixed $M \times$ overhead with no training; under selective operation only the abstaining fraction $1 - \phi(\tau)$ is escalated to deeper analysis, so the amortized cost is tuned by the same threshold that sets the risk-coverage operating point (§6).

5 Further Considerations

Each consideration below is handled by reusing a component of §4, not by adding a new one.

Unseen brands and the long tail. A page may impersonate a brand outside the canonical lexicon, where the value-normalization of §4.2 cannot resolve a logo-brand or brand-claim unit. The framework degrades gracefully rather than failing: the structural units (form-action, cert-org, redirect) require no lexicon and remain fully comparable and verifiable through (2)–(3), so GEA is still defined on the evidence the page does carry. The only effect is that an unresolvable brand unit drops out of the consensus; this costs no retraining and leaves the calibrator of §4.4 unchanged.

Bounded panels and budget. Every panel member is a forward pass, so the M -model cost is the binding budget constraint. The selective rule of §4.4 already bounds it: only the abstaining fraction is escalated, so the average cost is set by the same threshold τ that fixes the risk-coverage operating point. When budget is tight a two-stage variant runs a single model first and convenes the full panel only on its low-confidence pages; this reuses the abstention machinery unchanged and adds panel passes on an $O(1-\phi)$ fraction of traffic rather than on all of it.

New evidence types and channels. The schema $\mathcal{U}$ is open. A new signal – a unit drawn from the surrounding message-channel text when the URL arrives over chat, or from a community scam report – is admitted by registering one additional type and its verifier g ; the agreement and groundedness of (2)–(3) are defined over whatever types are present and need no change, and the calibrator re-fits on the same oracle. Extending coverage is therefore an $O(1)$ engineering step per signal, not a redesign of the score.

6 Experiments

We study five questions, each tied to a requirement of §3:

- **RQ1 – Selective detection (§6.2):** does grounded evidence-agreement give a better risk-coverage trade-off than label-level reliability?
- **RQ2 – Right-label, wrong-reason (§6.3):** does PHISH-PROOF flag the pages a panel labels correctly but for the wrong reason, which the baselines act on?
- **RQ3 – Component attribution (§6.4):** how much do evidence-level agreement, groundedness, diversity, and calibration each contribute?
- **RQ4 – Calibration and drift (§6.5):** is the trust score calibrated, and does it hold across a temporal split?
- **RQ5 – Cost (§6.6):** what does the panel cost under selective operation?

6.1 Setup

Tasks and datasets. We evaluate on PHISHSEL, a benchmark of roughly 4,000 webpages assembled over a six-month window: live phishing URLs drawn from public feeds [18, 20] and benign pages sampled from the Tranco top-sites list [11], each captured with its URL, DOM, rendered screenshot, TLS certificate, and redirect chain. Ground truth for evaluation comes from the delayed oracle – a phishing URL’s eventual blocklisting or take-down – recorded after a fixed lag, which also supplies the labels the calibrator of §4.4 consumes. The six-month span lets us hold out a later slice for the temporal study (RQ4).

Panel and baselines. The default panel is $M=3$ deliberately diverse detectors: a frontier text LLM, a second text LLM from a different family, and a vision-language model that reads the screenshot. We compare against six reliability baselines that share this panel or a member of it:

- *B1 Single-model confidence* [7]: the strongest single model with temperature-scaled softmax confidence.
- *B2 Verbalized confidence*: the same model asked to state its confidence.
- *B3 Self-consistency* [24]: one model, $k=5$ samples, agreement as the score.
- *B4 Majority vote* [22]: panel label agreement.
- *B5 Label-agreement* [4]: reliability-weighted label agreement over the panel.
- *B6 MultiPhishGuard* [25]: a multi-agent detector whose consolidated-rationale confidence is used as the score.

Tasks and metrics. All methods produce a per-page trust score and we evaluate them as selective detectors via (5). We report the area under the risk-coverage curve (*AURC*, selective risk integrated over coverage); selective accuracy and false-positive rate at 80% coverage (*SelAcc*₈₀, *FPR*₈₀); the coverage attainable at 99% selective accuracy (*Cov*₉₉); and the expected calibration error of the trust score (*ECE*). For RQ2 we also report the *flip rate*: the fraction of pages whose panel label changes under a small evidence-targeted perturbation (a logo morph [8] or a form-action cloak).

Statistical reporting. Every cell is the mean over five seeds with a 95% confidence interval from a paired bootstrap; we treat non-overlapping intervals as a significance test and state relative gains rather than raw point differences. The proposed method is **bold** and the best baseline underlined.

Reproducibility environment. The text models are queried through their providers’ APIs at temperature 0 with fixed seeds; the vision-language model and the perceptual logo verifier run on a single commodity GPU; the structural verifiers and the isotonic calibrator are CPU-only. We log every prompt, the raw model outputs, the cited units, and the verifier results, and release the schema, prompts, verifiers, calibration code, and the per-page score traces so that the risk-coverage curves can be reconstructed end to end.

6.2 Selective detection

Risk-coverage and deployment coverage. We first ask whether scoring agreement on evidence improves the core selective trade-off. Tab. 3 reports the headline comparison and Fig. 3 the full curves. PHISHPROOF lowers AURC by 35% relative to the strongest baseline

Table 3: RQ1. Label-free selective detection on PHISHSEL ($N \approx 4,000$ pages, 6-month span, $M=3$ diverse panel). AURC: area under the risk-coverage curve over coverage $[0, 1]$ (selective risk in points). SelAcc₈₀/FPR₈₀: selective accuracy and false-positive rate at 80% coverage. Cov₉₉: coverage attainable at 99% selective accuracy. ECE: expected calibration error. Cells are mean $\pm 95\%$ CI over 5 seeds. Best; best baseline. The Δ row is relative to the best baseline (B5), points for SelAcc.

Method	AURC↓	SelAcc ₈₀ ↑	FPR ₈₀ ↓	Cov ₉₉ ↑	ECE↓
B1 Single-model conf.	5.9 $\pm$ 0.4	95.8 $\pm$ 0.4	4.1 $\pm$ 0.5	0.38 $\pm$ 0.05	0.082 $\pm$ 0.009
B2 Verbalized conf.	6.8 $\pm$ 0.5	95.0 $\pm$ 0.5	4.9 $\pm$ 0.6	0.29 $\pm$ 0.05	0.110 $\pm$ 0.011
B3 Self-consistency ($k=5$)	5.6 $\pm$ 0.4	96.1 $\pm$ 0.4	3.8 $\pm$ 0.5	0.41 $\pm$ 0.05	0.074 $\pm$ 0.008
B4 Majority vote	5.3 $\pm$ 0.4	96.3 $\pm$ 0.4	3.6 $\pm$ 0.4	0.45 $\pm$ 0.05	0.069 $\pm$ 0.008
B5 Label-agreement	5.1 $\pm$ 0.4	96.5 $\pm$ 0.4	3.4 $\pm$ 0.4	0.48 $\pm$ 0.05	0.063 $\pm$ 0.007
B6 MultiPhishGuard	5.7 $\pm$ 0.5	96.0 $\pm$ 0.4	3.9 $\pm$ 0.5	0.42 $\pm$ 0.05	0.090 $\pm$ 0.010
PHISHPROOF (GEA)	3.3$\pm$0.3	98.4$\pm$0.3	1.6$\pm$0.3	0.71$\pm$0.04	0.031$\pm$0.005
Δ vs best baseline	-35%	+1.9	-53%	+48%	-51%




Figure 3: Risk-coverage on PHISHSEL (RQ1). Selective risk (100-selective accuracy, in points) vs. coverage; lower is better. All methods meet at the shared full-coverage panel error because PHISHPROOF changes the abstention order, not the label. The ringed point at 80% coverage reproduces SelAcc₈₀ of Tab. 3 (1.6 vs. 3.5). PHISHPROOF dominates at every coverage.

and roughly halves the false-positive rate at 80% coverage; most usefully for deployment, it raises the coverage attainable at 99% selective accuracy from 0.48 to 0.71, meaning an operator can act automatically on far more traffic at the same error budget. The mechanism is visible in Fig. 3: all methods meet at the shared full-coverage panel error – PHISHPROOF does not change the majority label, only the order in which pages are trusted – but its curve falls fastest as coverage tightens, because grounded evidence-agreement ranks the safe pages above the fragile ones more sharply than label concordance does. The label-aggregation baselines (B4, B5) edge out single-model confidence but remain well above PHISHPROOF, confirming that the gain comes from *where* agreement is measured, not merely from using several models.

Table 4: RQ2. Anatomy of the *right-label, wrong-reason* (RLWR) failure, over pages the $M=3$ panel labels *correctly*. Each region is defined by PHISHPROOF’s two factors (evidence agreement A , groundedness G). “B1 conf.”: mean single-model confidence; “Unanim.”: fraction with a unanimous panel label (B4); “Abstain”: fraction PHISHPROOF withholds; “Flip”: label-flip rate under a small evidence-targeted perturbation (logo morph / form-action cloak). Baselines act confidently on RLWR and ungrounded pages, which are precisely the fragile ones; PHISHPROOF abstains.

Region (of correct labels)	Share	B1 conf.	Unanim.	Abstain	Flip↓
High-GEA ($A \uparrow, G \uparrow$)	72%	0.93	0.89	4%	6.3%
RLWR ($A \downarrow$ on evidence)	11%	0.91	0.64	78%	41.6%
Ungrounded ($G \downarrow$, shared halluc.)	6%	0.88	0.71	83%	47.9%
Low label-agreement (uncertain)	11%	0.74	0.18	81%	33.2%

6.3 Right-label, wrong-reason

Partitioning correct verdicts by grounded agreement. PHISHPROOF’s premise is that some correctly labeled pages are untrustworthy because the panel agrees on the label but not on grounded evidence. We test this by partitioning the correctly labeled pages with PHISHPROOF’s two factors (Tab. 4). About 11% fall in the right-label, wrong-reason region and a further 6% rest on ungrounded (shared-hallucination) evidence. These regions are not benign: under a small evidence-targeted perturbation their labels flip 41.6% and 47.9% of the time, against 6.3% for high-GEA pages – they are exactly the pages a minor change to the cited cue overturns. Yet the baselines act on them confidently, with mean single-model confidence 0.91 and a unanimous panel label on 64% of the RLWR region; label-level reliability cannot see the problem because the labels do agree. PHISHPROOF abstains on 78–83% of these pages while acting on the stable high-GEA region, and Fig. 4(b) shows the same effect continuously: adversarial fragility falls monotonically from the lowest to the highest PHISHPROOF decile.

6.4 Component attribution

Per-component contribution. ?? removes one component at a time. Replacing evidence-level agreement with label agreement is the most damaging change: AURC rises to 5.0, essentially reproducing the B5 baseline of Tab. 3 and confirming that the contribution is the evidence channel, not the panel. Dropping groundedness admits shared hallucinations and costs the next most, and collapsing the panel to three same-family variants – removing diversity while keeping the count – erases much of the remaining advantage, which is why diversity, not panel size, is the design variable. Removing oracle calibration leaves ranking nearly intact but degrades the trust score’s ECE from 0.031 to 0.071, i.e. it mostly hurts the *reported* probability rather than the *order*. The groundedness verifiers are themselves accurate – the structural checks are exact by construction, and the perceptual logo verifier attains 0.96/0.93 precision/recall against a gold logo-equivalence set – so G seldom discards a true unit.




Figure 4: Diagnostics. (a) **Reliability:** PHISHPROOF’s calibrated trust tracks the diagonal, whereas single-model confidence is over-confident; the annotated ECE matches Tab. 3. (b) **Label-flip rate** under a small evidence-targeted perturbation, by PHISHPROOF decile: the low-GEA (RLWR) region is adversarially fragile while high-GEA pages are stable, reproducing the flip rates of Tab. 4.

6.5 Calibration and drift

Reliability and temporal robustness. A selective detector is only as useful as its operating point, so we check whether the calibrated score is faithful and whether it survives drift. Fig. 4(a) plots reliability: PHISHPROOF’s calibrated trust tracks the diagonal, whereas single-model confidence sits below it – systematically over-confident – consistent with the ECE column of Tab. 3. To test drift we calibrate on the first three months and evaluate on the last three, a temporal protocol in the spirit of drift-aware malware evaluation [19] (??). Both methods degrade as campaigns evolve, but PHISHPROOF retains its margin (AURC 3.3→3.7 versus 5.1→5.6 for B5), because grounded evidence-agreement depends on page-level cues that move more slowly than the surface features a label classifier keys on.

6.6 Cost under selective operation

Overhead and selective escalation. PHISHPROOF adds a fixed $M=3\times$ model overhead over a single detector and no training; the structural verifiers are constant-time page lookups and the single perceptual check is one embedding comparison, so per-page verification is negligible beside the model calls. Selective operation makes even the model cost elastic: at the 80% operating point only the abstaining fifth of traffic is escalated to deeper analysis, and the two-stage variant of §5 runs the full panel only on a single model’s low-confidence pages, cutting average panel invocations toward that same fraction. Cost is thus governed by the very threshold that sets the risk–coverage point, and trades smoothly against it.

7 Conclusion

Deployed phishing detectors are trusted at the label and explained, if at all, after the fact, even though the action a verdict triggers needs a reason an operator can check. We have argued that reliability should be read from the evidence rather than the label, and we instantiated this as PHISHPROOF: a multi-model detector whose Grounded Evidence-Agreement score trusts a verdict in proportion

to how much independent models agree on evidence that the page itself confirms, returns that agreed evidence as the explanation, and abstains otherwise. Three findings stand out. The right-label, wrong-reason page – correctly labeled but resting on disagreed or hallucinated evidence – is real, common, and the most adversarially fragile region of the input space, yet label confidence and majority voting act on it. Measuring agreement on evidence rather than on labels is what carries the improvement, as removing that single component returns the detector to a label-agreement baseline. And because the score is grounded in page-level cues, it stays calibrated and keeps its margin as campaigns drift. Trust and explanation, in short, can be decoupled from raw accuracy and earned per verdict.

The detector opens several directions that extend it without re-designing the score. A verifier is sound for what it checks but a page can carry true-yet misleading evidence, so strengthening verifiers toward semantic plausibility would tighten the grounding guarantee. The evidence schema and the panel are fixed here; learning new evidence types and inferring how diverse a panel actually is would let the method adapt to new impersonation tactics and to opaque third-party model pools. And because the calibrating oracle arrives late, the hardest regime remains the fully cloaked, novel page on which the detector can today only abstain; narrowing that gap with faster, page-internal grounding is the natural next step. Each direction reuses the same evidence-agreement machinery rather than replacing it.

References

- [1] Sahar Abdelnabi, Katharina Krombholz, and Mario Fritz. 2020. VisualPhishNet: Zero-Day Phishing Website Detection by Visual Similarity. In *ACM SIGSAC Conference on Computer and Communications Security (CCS)*.
- [2] Christina Baek, Yiding Jiang, Aditi Raghunathan, and J.Žico Kolter. 2022. Agreement-on-the-Line: Predicting the Performance of Neural Networks under Distribution Shift. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:2206.13089.
- [3] Xuan Chen, Hao Liu, et al. 2026. MemoPhishAgent: Memory-Augmented Multi-Modal LLM Agent for Phishing URL Detection. arXiv:2602.21394 arXiv:2602.21394; under review at ICLR 2026.
- [4] A.P. Dawid and A.M. Skene. 1979. Maximum Likelihood Estimation of Observer Error-Rates Using the EM Algorithm. *Journal of the Royal Statistical Society*.

- Series C (Applied Statistics)* 28, 1 (1979), 20–28.
- [5] Daniyal Ganiuly and Assel Smayil. 2025. Trustworthiness Calibration Framework for Phishing Email Detection Using Large Language Models. arXiv:2511.04728 arXiv:2511.04728.
 - [6] Yonatan Geifman and Ran El-Yaniv. 2017. Selective Classification for Deep Neural Networks. In *Advances in Neural Information Processing Systems (NeurIPS)*. arXiv:1705.08500.
 - [7] Chuan Guo, Geoff Pleiss, Yu Sun, and Kilian Q. Weinberger. 2017. On Calibration of Modern Neural Networks. In *International Conference on Machine Learning (ICML)*. arXiv:1706.04599.
 - [8] Qingying Hao, Nirav Diwan, Ying Yuan, Giovanni Apruzzese, Mauro Conti, and Gang Wang. 2024. It Doesn't Look Like Anything to Me: Using Diffusion Model to Subvert Visual Phishing Detectors. In *USENIX Security Symposium*.
 - [9] Sourena Khanzadeh. 2026. Project Ariadne: A Structural Causal Framework for Auditing Faithfulness in LLM Agents. arXiv:2601.02314 arXiv:2601.02314.
 - [10] Kyungha Kim, Sangyun Lee, Kung-Hsiang Huang, et al. 2024. Can LLMs Produce Faithful Explanations For Fact-checking? Towards Faithful Explainable Fact-Checking via Multi-Agent Debate. arXiv:2402.07401 arXiv:2402.07401.
 - [11] Victor Le Pochat, Tom Van Goethem, Samaneh Tajalizadehkhoob, Maciej Korczyński, and Wouter Joosen. 2019. Tranco: A Research-Oriented Top Sites Ranking Hardened against Manipulation. In *Network and Distributed System Security Symposium (NDSS)*. arXiv:1806.01156.
 - [12] Wenhao Li, Selvakumar Manickam, Yung-Wey Chong, and Shankar Karuppayah. 2025. PhishDebate: An LLM-Based Multi-Agent Framework for Phishing Website Detection. arXiv:2506.15656 arXiv:2506.15656.
 - [13] Yuexin Li, Chengyu Huang, Shumin Deng, Mei Lin Lock, Tri Cao, Nay Oo, Hoon Wei Lim, and Bryan Hooi. 2024. KnowPhish: Large Language Models Meet Multimodal Knowledge Graphs for Enhancing Reference-Based Phishing Detection. In *USENIX Security Symposium*. arXiv:2403.02253.
 - [14] Yun Lin, Ruofan Liu, Dinil Mon Divakaran, Jun Yang Ng, Qing Zhou Chan, Yiwen Lu, Yuxuan Si, Fan Zhang, and Jin Song Dong. 2021. Phishpedia: A Hybrid Deep Learning Based Approach to Visually Identify Phishing Webpages. In *USENIX Security Symposium*.
 - [15] Ruofan Liu, Yun Lin, Xiwen Teoh, Gongshen Liu, Zhiyong Huang, and Jin Song Dong. 2024. Less Defined Knowledge and More True Alarms: Reference-based Phishing Detection without a Pre-defined Reference List. In *USENIX Security Symposium*. 523–540.
 - [16] Ruofan Liu, Yun Lin, Xianglin Yang, Siang Hwee Ng, Dinil Mon Divakaran, and Jin Song Dong. 2022. Inferring Phishing Intention via Webpage Appearance and Dynamics: A Deep Vision Based Approach. In *USENIX Security Symposium*.
 - [17] Ruofan Liu, Yun Lin, Yifan Zhang, Penn Han Lee, and Jin Song Dong. 2023. Knowledge Expansion and Counterfactual Interaction for Reference-Based Phishing Detection. In *USENIX Security Symposium*.
 - [18] OpenPhish. [n. d.]. OpenPhish: Phishing Intelligence Feed. <https://openphish.com>. Accessed during protocol design.
 - [19] Feargus Pendlebury, Fabio Pierazzi, Roberto Jordaney, Johannes Kinder, and Lorenzo Cavallaro. 2019. TESSERACT: Eliminating Experimental Bias in Malware Classification across Space and Time. In *USENIX Security Symposium*.
 - [20] PhishTank. [n. d.]. PhishTank: Community-Based Anti-Phishing Feed. <https://phishtank.org>. Accessed during protocol design.
 - [21] Asadullah Safi and Satwinder Singh. 2023. A Systematic Literature Review on Phishing Website Detection Techniques. *Journal of King Saud University – Computer and Information Sciences* 35, 2 (2023), 590–611.
 - [22] Fouad Trad and Ali Chehab. 2024. To Ensemble or Not: Assessing Majority Voting Strategies for Phishing Detection with Large Language Models. arXiv:2412.00166 arXiv:2412.00166.
 - [23] Vladimir Vovk, Alexander Gammerman, and Glenn Shafer. 2005. *Algorithmic Learning in a Random World*. Springer.
 - [24] Xuezhi Wang, Jason Wei, Dale Schuurmans, Quoc Le, Ed Chi, Sharan Narang, Aakanksha Chowdhery, and Denny Zhou. 2023. Self-Consistency Improves Chain of Thought Reasoning in Language Models. In *International Conference on Learning Representations (ICLR)*. arXiv:2203.11171.
 - [25] Yinuo Xue, Eric Spero, Meng Wai Woo, Wei Gao, and Giovanni Russello. 2025. MultiPhishGuard: An Explainable and Adaptive Multi-Agent LLM System for Phishing Email Detection. arXiv:2505.23803 arXiv:2505.23803.
 - [26] Xu Yang et al. 2025. Phishing Website Detection on Imbalanced Datasets Based on Diffusion Model and Multi-Head Self-Attention. In *Intl. Conf. on Communications, Signal Processing and Information (CSIP)*. doi:10.1145/3744668.3744669